

# NB\_6\_custom\_Loss

November 22, 2025

```
[9]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.metrics import precision_score, recall_score, f1_score, accuracy_score
import torch
import torch.nn as nn
import torch.nn.functional as F

# Set random seeds for reproducibility
torch.manual_seed(0)
np.random.seed(0)

print("Imports OK ")
```

Imports OK

```
[10]: # 1. Load Data
filename = "/Users/gasper/Documents/PSL/year_1/semster_1/
˓→Practical-Machine-Learning/NB6 - Gradient Descent/Fraud_detection.csv"
df = pd.read_csv(filename)
df_clean = df.dropna()

print(f"Dataset shape: {df_clean.shape}")

# 2. Identify Amount Column
# We need to extract the Amount BEFORE splitting to ensure we don't lose track
˓→of it
if "Amount" in df_clean.columns:
    AMOUNT_COL = "Amount"
elif "amount" in df_clean.columns:
    AMOUNT_COL = "amount"
else:
    # Fallback: try to find the first numeric column that isn't the target
    cols = df_clean.select_dtypes(include=[np.number]).columns.tolist()
    if "target" in cols:
```

```

        cols.remove("target")
AMOUNT_COL = cols[0]

print(f"Using '{AMOUNT_COL}' as the transaction amount column.")

# 3. Split Data (X, y, and Amounts)
# We split X and y as usual, but we also need to split the Amounts array
# so it matches the X/y indices exactly.
X = df_clean.drop(columns=["target"]).values
y = df_clean["target"].values
amounts = df_clean[AMOUNT_COL].values

# Train/Temp Split
X_tr, X_tmp, y_tr, y_tmp, amt_tr, amt_tmp = train_test_split(
    X, y, amounts, test_size=0.4, random_state=42, stratify=y
)

# Val/Test Split
X_va, X_te, y_va, y_te, amt_va, amt_te = train_test_split(
    X_tmp, y_tmp, amt_tmp, test_size=0.5, random_state=43, stratify=y_tmp
)

# 4. Convert to PyTorch Tensors
# Note: We need amounts as float32 for the loss calculation
X_train_t = torch.tensor(X_tr, dtype=torch.float32)
y_train_t = torch.tensor(y_tr, dtype=torch.float32) # Float needed for custom loss
amt_train_t = torch.tensor(amt_tr, dtype=torch.float32)

X_val_t = torch.tensor(X_va, dtype=torch.float32)
y_val_t = torch.tensor(y_va, dtype=torch.float32)
amt_val_t = torch.tensor(amt_va, dtype=torch.float32)

X_test_t = torch.tensor(X_te, dtype=torch.float32)
y_test_t = torch.tensor(y_te, dtype=torch.float32)
amt_test_t = torch.tensor(amt_te, dtype=torch.float32)

print(f"Train set: {X_train_t.shape}")
print(f"Val set: {X_val_t.shape}")
print(f"Test set: {X_test_t.shape}")

```

Dataset shape: (164492, 30)  
Using 'Amount' as the transaction amount column.  
Train set: torch.Size([98695, 29])  
Val set: torch.Size([32898, 29])  
Test set: torch.Size([32899, 29])

```
[11]: class ExpectedCostLoss(nn.Module):
    """
    Custom loss function that minimizes total business cost.
    """

    def __init__(self, abandonment_limit=0.1, user_lazy_factor=0.2, ↴
                 penalty_per_percent=10000):
        super().__init__()
        self.limit = abandonment_limit
        self.lazy_factor = user_lazy_factor
        # 10,000 EUR penalty for every 1% (0.01) means we multiply by 1,000,000
        # Example: 0.11 - 0.10 = 0.01 * 1,000,000 = 10,000
        self.penalty_scaler = penalty_per_percent * 100

    def forward(self, logits, targets, amounts):
        """
        logits: (batch_size, 1) Raw outputs
        targets: (batch_size, ) 0 for Legit, 1 for Fraud
        amounts: (batch_size, ) Transaction values in EUR
        """
        # 1. Get Probabilities
        probs = torch.sigmoid(logits).squeeze()

        # 2. Authentication Cost
        # We pay ~1 EUR for every high probability assignment
        auth_cost = probs * 1.0

        # 3. Missed Fraud Cost
        # If target=1 and prob=0.1, we miss 0.9 * Amount
        missed_fraud_cost = targets * (1 - probs) * amounts

        # 4. Abandonment Penalty (Global Batch Constraint)
        avg_auth_rate = probs.mean()
        current_abandon_rate = avg_auth_rate * self.lazy_factor

        # ReLU handles the "cliff": if rate < limit, penalty is 0
        excess_abandonment = F.relu(current_abandon_rate - self.limit)
        abandonment_penalty = excess_abandonment * self.penalty_scaler

        # 5. Total Loss
        # We average per-transaction costs and add the global penalty
        total_loss = (auth_cost + missed_fraud_cost).mean() + ↴
                     abandonment_penalty

    return total_loss
```

```
[12]: class CostOptimizedNN(nn.Module):
    def __init__(self, input_dim, hidden_dim=32):
        super().__init__()
        self.net = nn.Sequential(
            nn.Linear(input_dim, hidden_dim),
            nn.ReLU(),
            nn.Linear(hidden_dim, 1), # Output 1 scalar (logit)
        )

    def forward(self, x):
        return self.net(x)

# Initialize
model = CostOptimizedNN(input_dim=X_train_t.shape[1])
criterion = ExpectedCostLoss()
optimizer = torch.optim.Adam(model.parameters(), lr=0.001)

epochs = 100
batch_size = 64
N = X_train_t.shape[0]

# Track losses
train_losses = []
val_losses = []

print("Starting Cost-Sensitive Training...")

for epoch in range(epochs):
    model.train()
    perm = torch.randperm(N)
    total_loss = 0.0

    for start in range(0, N, batch_size):
        idx = perm[start : start + batch_size]
        xb = X_train_t[idx]
        yb = y_train_t[idx]
        ab = amt_train_t[idx] # Batch amounts

        # Forward
        logits = model(xb)
        loss = criterion(logits, yb, ab)

        # Backward
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()
```

```

    total_loss += loss.item()

    # Calculate average training loss for the epoch
    avg_train_loss = total_loss / N
    train_losses.append(avg_train_loss)

    # Validation Check
    model.eval()
    with torch.no_grad():
        val_logits = model(X_val_t)
        val_loss = criterion(val_logits, y_val_t, amt_val_t).item()
        val_losses.append(val_loss)

    if epoch % 10 == 0:
        print(
            f"Epoch {epoch}: Train Loss: {avg_train_loss:.4f} | Val Cost Score: {val_loss:.2f}"
        )

print("Training Complete.")

```

Starting Cost-Sensitive Training...

Epoch 0: Train Loss: 4.6067 | Val Cost Score: 2.76

Epoch 10: Train Loss: 0.0365 | Val Cost Score: 2.42

Epoch 20: Train Loss: 0.0218 | Val Cost Score: 2.29

Epoch 30: Train Loss: 0.0341 | Val Cost Score: 2.30

Epoch 40: Train Loss: 0.0034 | Val Cost Score: 0.23

Epoch 50: Train Loss: 0.0039 | Val Cost Score: 0.31

Epoch 60: Train Loss: 0.0039 | Val Cost Score: 0.31

Epoch 70: Train Loss: 0.0057 | Val Cost Score: 0.24

Epoch 80: Train Loss: 0.0038 | Val Cost Score: 0.24

Epoch 90: Train Loss: 0.0036 | Val Cost Score: 0.23

Training Complete.

```
[ ]: # Plot Training and Validation Loss
plt.figure(figsize=(10, 6))

plt.plot(train_losses, label="Training Loss", linewidth=2)
plt.plot(val_losses, label="Validation Loss", linewidth=2)
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.title("Training vs Validation Loss")
plt.legend()
plt.grid(True, alpha=0.3)

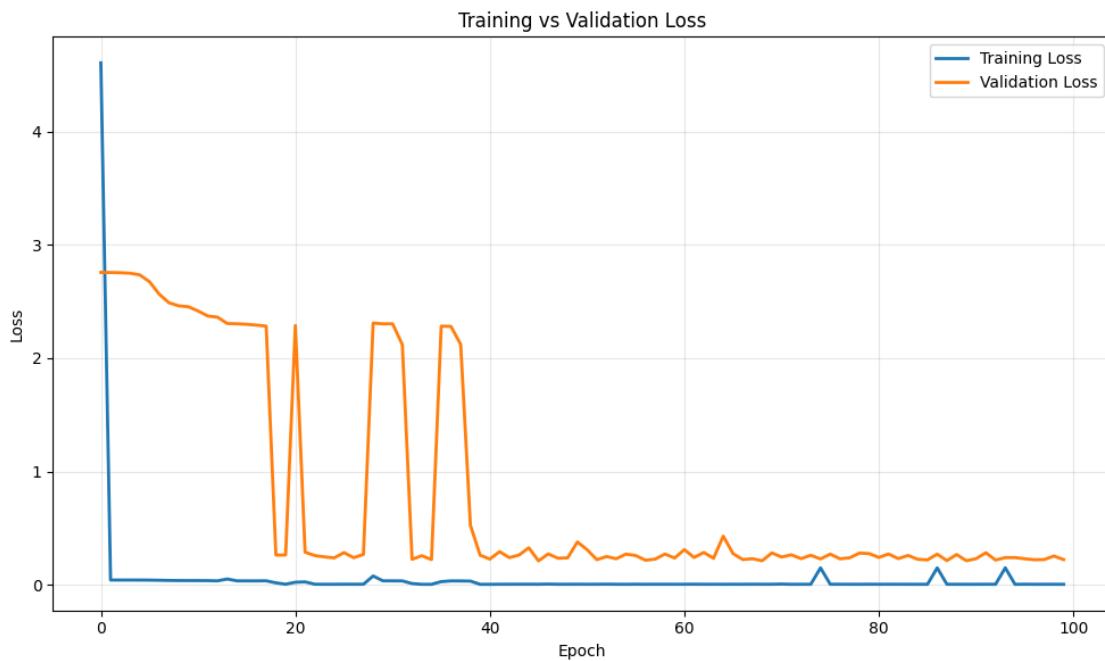
plt.tight_layout()
```

```

plt.show()

print(f"\nFinal Training Loss: {train_losses[-1]:.4f}")
print(f"Final Validation Loss: {val_losses[-1]:.2f}")

```



Final Training Loss: 0.0034

Final Validation Loss: 0.22

```

[18]: def calculate_bank_invoice(model, X_tensor, y_true, amounts, threshold=0.5):
    model.eval()
    with torch.no_grad():
        logits = model(X_tensor)
        probs = torch.sigmoid(logits).squeeze().numpy()

    # Make hard decisions based on threshold
    preds = (probs >= threshold).astype(int)

    total_transactions = len(y_true)

    # 1. Authentication Costs
    auth_count = (preds == 1).sum()
    auth_fee = auth_count * 1.0

    # 2. Missed Fraud Costs
    # We let it go (pred=0) but it was fraud (y=1)

```

```

missed_indices = np.where((preds == 0) & (y_true == 1))[0]
missed_fraud_cost = amounts[missed_indices].sum()

# 3. Abandonment Penalty
abandon_rate = (auth_count * 0.2) / total_transactions
excess = max(0, abandon_rate - 0.10)
# 10k for every 1% (0.01)
abandon_penalty = 10000 * (excess * 100)

total_cost = auth_fee + missed_fraud_cost + abandon_penalty

print(f"--- FINAL BANK INVOICE ---")
print(f"Auth Requested: {auth_count} ({auth_fee:.2f})")
print(f"Missed Fraud Cost: {missed_fraud_cost:.2f}")
print(f"Abandonment Rate: {abandon_rate:.2%} (Penalty: {abandon_penalty:.2f})")
print(f"-----")
print(f"TOTAL COST: {total_cost:.2f}")

return total_cost, preds

```

*# Run Evaluation on Test Set*

```

total_cost, final_preds = calculate_bank_invoice(
    model,
    X_test_t,
    y_te, # Use the numpy array for y_true
    amt_te, # Use the numpy array for amounts
    threshold=0.5,
)

```

```

--- FINAL BANK INVOICE ---
Auth Requested: 4511 (€4511.00)
Missed Fraud Cost: €2,882.21
Abandonment Rate: 2.74% (Penalty: €0.00)
-----
TOTAL COST: €7,393.21

```

[19]:

```

# Get probabilities for plotting
with torch.no_grad():
    probs = torch.sigmoid(model(X_test_t)).squeeze().numpy()

plt.figure(figsize=(10, 6))

# Scatter plot: X-axis = Amount, Y-axis = Predicted Probability
# Color = Actual Class (Red=Fraud, Blue=Legit)

```

```

subset = np.random.choice(len(amt_te), size=min(1000, len(amt_te)),  

    ↪replace=False)

plt.scatter(  

    amt_te[subset][y_te[subset] == 0],  

    probs[subset][y_te[subset] == 0],  

    alpha=0.3,  

    label="Legit",  

    color="blue",  

    s=10,  

)  

plt.scatter(  

    amt_te[subset][y_te[subset] == 1],  

    probs[subset][y_te[subset] == 1],  

    alpha=0.6,  

    label="Fraud",  

    color="red",  

    s=20,  

)  

plt.axhline(0.5, color="gray", linestyle="--", label="Threshold")  

plt.xlabel("Transaction Amount (€)")  

plt.ylabel("Model Probability of Fraud")  

plt.title("Cost-Sensitive Logic: Does the model catch high-value frauds?")  

plt.legend()  

plt.show()

```

